

Deep Face Detection Model - Milestone #3

Brayan Leonardo Gil Guevara (C0902422)

Eduardo Williams (C0896405)

Marzieh Mohammadi Kokaneh (C089839)

Rohit Kumar (C0895100)

Saurabh Laltaprasad Gangwar (C0894380)

Lambton College

Big Data Analytics

BDM 3035 - Big Data Capstone Project 01

PhD. Meysam Effati

Mississauga, Ontario

July 1st, 2024

Contents

Introduction.....	3
Progress Report.....	4
Building Model.....	11
Defining the FaceTracker Model.....	15
Compiling and Training the Model.....	18
Making Predictions on the Test Dataset	18
Results	20
Next Steps	21
Faced Challenges.....	22
Lesson Learned	24
Conclusion	26

Introduction

The face recognition industry has been crucial for security purposes over the years, however, once COVID hit in late 2019, more requirements for this technology have been raised. This document explores the use of deep learning for facial recognition. The Idea comes from creating a local dataset having different photos of people on a team and training the model to detect the identity of each person. We will discuss how to collect and prepare data, as well as develop models for this combined task. The target for this project is security companies for letting people enter the office or not! For this specific document we are focusing on Data Splitting Process and Data Augmentation as we are getting ready for NL Model Training and Performance Analysis in future steps.

Progress Report

Our next step for this milestone was to build and train the model.

We had already distributed all the jpg files from the images folder as follows:

Train folder → 70%

Test folder → 15%

Valid folder → 15%

Importing necessary modules:

We imported the following models:

```
import pandas as pd
import numpy as np
import os
import cv2
import json
from matplotlib import pyplot as plt

# Importing layers and base network
import tensorflow as tf
from tensorflow.keras.models import load_model, Model
from tensorflow.keras.layers import Input, Dense, GlobalMaxPooling2D
from tensorflow.keras.applications import VGG16
```

Figure 1. Importing the layers and base network

Load_model

- Description: The load_model function is used to load a previously saved Keras model from a file. This includes the model architecture, weights, and optimizer state.
- Usage: This function is commonly used when you need to deploy a trained model, continue training, or evaluate it on new data without having to redefine the model structure.

Model

- Description: The Model class is used to create a Keras model with more flexibility compared to the Sequential API. It allows the creation of complex models with multiple inputs and outputs or shared layers.
- Usage: This class is ideal for defining models that do not fit into a simple stack of layers, such as models with multiple branches or skip connections.

Input

- Description: The Input class is used to instantiate a Keras tensor, which acts as a placeholder for the input data to the model. It is essential in the functional API to define the input shape of the model.
- Usage: This is used as the first layer in a functional model, defining the shape and data type of the input.

Dense

- Description: The Dense layer is a fully connected neural network layer. It is used to connect every neuron in the previous layer to every neuron in the current layer.
- Usage: This layer is commonly used in feedforward neural networks, especially in the final classification or regression layers.

GlobalMaxPooling2D

- Description: The GlobalMaxPooling2D layer performs global max pooling operation on the spatial dimensions (height and width) of the input tensor. This operation reduces each feature map to a single value by taking the maximum value of all elements in each feature map.

- Usage: This layer is typically used to reduce the spatial dimensions of the data while preserving the most important features, often used in CNN architectures before the fully connected layers.

VGG16

- Description: Convolutional neural network, also known as ConvNet, is a type of artificial neural network. A convolutional neural network consists of an input layer, an output layer, and multiple hidden layers.

VGG16 is a type of CNN (Convolutional Neural Network) and is considered one of the best computer models to date. The design scaled the mesh and added depth using a small convolutional filter design, which represents a significant improvement over current technology.

By pushing the depth to a 16x19 weight layer, they give about 138 parameters. It classifies 1000 different groups with 92.7% accuracy. It is one of the popular algorithms for image classification and is easy to use with transfer learning.

VGG16 has a total of 21 layers, including 13 convolutional layers, 5 max-pooling layers, and 3 thickness layers, but only 16 layers are used to learn the process parameter. 244 and 3 RGB channels




Figure 2: VGG16 architecture

Usage: The most special thing about VGG16 is that they do not use too many hyperparameters, with step 1 they focus on the convolutional process with a 3x3 filter, and with step 2 they always use 2x2. Same padding and maxpool set to filter. Transformation 3 has 256 filters, Transformation 4 and Transformation 5 have 512 filters. Each performs a 1000-way ILSVRC classification and therefore has 1000 channels (one channel per class). The last layer is the soft-max layer.

Dataset Augmentation Paths:

We defined the directory path for the augmented dataset which included training, testing and validation images and labels.

This step helps us organize the data and makes it accessible for processing.

```
# defining paths
diraug_train_img = "data_aug/train/images/"
diraug_train_lab = "data_aug/train/labels/"
diraug_test_img = "data_aug/test/images/"
diraug_test_lab = "data_aug/test/labels/"
diraug_valid_img = "data_aug/valid/images/"
diraug_valid_lab = "data_aug/valid/labels/"
```

Figure 3: Defining Paths

Image and Label Loading Functions

Next, we developed the functions which will load the image and its label

The “Load_image” function will load, read and decode the JPEG images while the “Load_labels” function will load and read the JSON files which contain the information about the image.

```
# defining Load_image function

def load_image(x):
    byte_img = tf.io.read_file(x)
    img = tf.io.decode_jpeg(byte_img)
    return img

# defining Label loading function

def load_labels(label_path):
    with open(label_path.numpy(), 'r', encoding = "utf-8") as f:
        label = json.load(f)

    return [label['class']], label['bbox']
```

Figure 4: Defining Load_image and Label function

Loading and Preprocessing Datasets:

We loaded the trained and augmented image to tensorflow dataset.

Here we resized the images to a consistent shape which is 120x120 pixels. This ensured that we had uniform data across.

Next, we divided the pixel values by dividing it with 255. This will help us scale the data between 0 and 1. We will have an improved performance of neural network models.

We did it for test, train and validation augmented image.

```
# Load train augmented images to tensorflow dataset
train_images = tf.data.Dataset.list_files(diraug_train_img + '*.jpg', shuffle=False)
train_images = train_images.map(load_image)
train_images = train_images.map(lambda x: tf.image.resize(x, (120,120)))
train_images = train_images.map(lambda x: x/255)
```

```
# Load test augmented images to tensorflow dataset
test_images = tf.data.Dataset.list_files(diraug_test_img + '*.jpg', shuffle=False)
test_images = test_images.map(load_image)
test_images = test_images.map(lambda x: tf.image.resize(x, (120,120)))
test_images = test_images.map(lambda x: x/255)
```

```
# Load valid augmented images to tensorflow dataset
val_images = tf.data.Dataset.list_files(diraug_valid_img + '*.jpg', shuffle=False)
val_images = val_images.map(load_image)
val_images = val_images.map(lambda x: tf.image.resize(x, (120,120)))
val_images = val_images.map(lambda x: x/255)
```


Figure 5: Loading augmented images to the tensorflow dataset

Later, we loaded the test, train, and validation labels to the tensor flow dataset.

```
# Load train Labels to Tensorflow Dataset
train_labels = tf.data.Dataset.list_files(diraug_train_lab + '*.json', shuffle=False)
train_labels = train_labels.map(lambda x: tf.py_function(load_labels, [x], [tf.uint8, tf.float16]))

# Load test Labels to Tensorflow Dataset
test_labels = tf.data.Dataset.list_files(diraug_test_lab + '*.json', shuffle=False)
test_labels = test_labels.map(lambda x: tf.py_function(load_labels, [x], [tf.uint8, tf.float16]))

# Load valid Labels to Tensorflow Dataset
val_labels = tf.data.Dataset.list_files(diraug_valid_lab + '*.json', shuffle=False)
val_labels = val_labels.map(lambda x: tf.py_function(load_labels, [x], [tf.uint8, tf.float16]))
```

Figure 6: Loading the labels to the tensorflow dataset

Combining and Preparing Final Datasets:

We combined the images and the labels of test, train and validation dataset.

This involved creating the pairs of images and corresponding labels. We prepared the dataset for model training by:

- Shuffling the data so that the model sees random values during the training
- We created mini batches so that it enables efficient data for training.
- We pre-loaded the data to enhance input pipeline efficiency and alleviate I/O constraints.

```
# creating final train dataset by combining labels and images samples
train = tf.data.Dataset.zip((train_images, train_labels))
train = train.shuffle(5000)
train = train.batch(8)
train = train.prefetch(4)

# creating final test dataset by combining labels and images samples
test = tf.data.Dataset.zip((test_images, test_labels))
test = test.shuffle(1300)
test = test.batch(8)
test = test.prefetch(4)

# creating final validation dataset by combining labels and images samples
val = tf.data.Dataset.zip((val_images, val_labels))
val = val.shuffle(1000)
val = val.batch(8)
val = val.prefetch(4)
```

Figure 7: Combining and preparing the final dataset.

Loop to Draw Bounding Boxes on Images

Here we extracted the image, and its co-ordinates from the “res” for each index idx ranging from 0 to 1

We have drawn the bounding box on the image using cv2.rectangle function provided by the OpenCV. The coordinates are adjusted to image size by multiplying with [120, 120] and then changed to integer type.

Next, we displayed the image featuring the bounding box which is shown as the output.

```
# drawing a sample
fig, ax = plt.subplots(ncols = 4, figsize = (20,20))

for idx in range(4):
    sample_image = res[0][idx].copy()
    sample_coord = res[1][1][idx]

    cv2.rectangle(sample_image,
                  tuple(np.multiply(sample_coord[:2], [120,120]).astype(int)),
                  tuple(np.multiply(sample_coord[2:], [120,120]).astype(int)),
                  (255,0,0), 2)

    #ax[idx].set_title(str(res[1][2][idx]))
    ax[idx].imshow(sample_image)
```

Figure 9: Loop to Draw Bounding Boxes on Images

Output:



Figure 10: Output

Building Model

VGG16 Base Model:

The VGG16 model, trained on ImageNet beforehand, serves as the primary feature extractor.

The higher layers are eliminated to keep only the convolutional layers.

The parameter “include_top=False” makes sure that the top fully connected layers of the VGG16 model are not included, enabling the extracted features to be utilized in the custom classification and localization sections.

```
# 8. Building a deep Learning using the functional API
# 8.2 Download VGG16
vgg = VGG16(include_top=False)
```

Total params: 14,714,688 (56.13 MB)

Trainable params: 14,714,688 (56.13 MB)

Non-trainable params: 0 (0.00 B)

Figure 11: Building Deep Learning

Building the Custom Model

A custom model is built using the Keras Functional API. The model is made up of two branches: one for classification and the other for localization.

Defines the input layer with a size of 120x120x3.

Features are extracted by running inputs through the VGG16 model.

Sets up two separate departments:

- Classification Branch: Employing global max pooling, two dense layers with ReLU and softmax activations.
- Localization Branch: Employing global max pooling, incorporating two dense layers with ReLU and sigmoid activation functions.

Combines the divisions into a single cohesive model.

```
# Build instance of Network
def build_model():
    input_layer = Input(shape=(120,120,3))

    vgg = VGG16(include_top=False)(input_layer)

    # classification model
    f1 = GlobalMaxPooling2D()(vgg)
    class1 = Dense(2048, activation='relu')(f1)
    class2 = Dense(3, activation='softmax')(class1)

    # regression model / bounding box model
    f2 = GlobalMaxPooling2D()(vgg)
    regress1 = Dense(2048, activation='relu')(f2)
    regress2 = Dense(4, activation='sigmoid')(regress1)

    facetracker = Model(inputs=input_layer, outputs=[class2, regress2])

    return facetracker
```

Figure 12: Building instance of Network

Model Summary

The model summary provides a summary of the layers and parameters.

The summary feature shows the configuration of the customized model,

specifying the kind and dimensions of each layer, along with the number of parameters, offering a glimpse into the model's intricacy and scale.

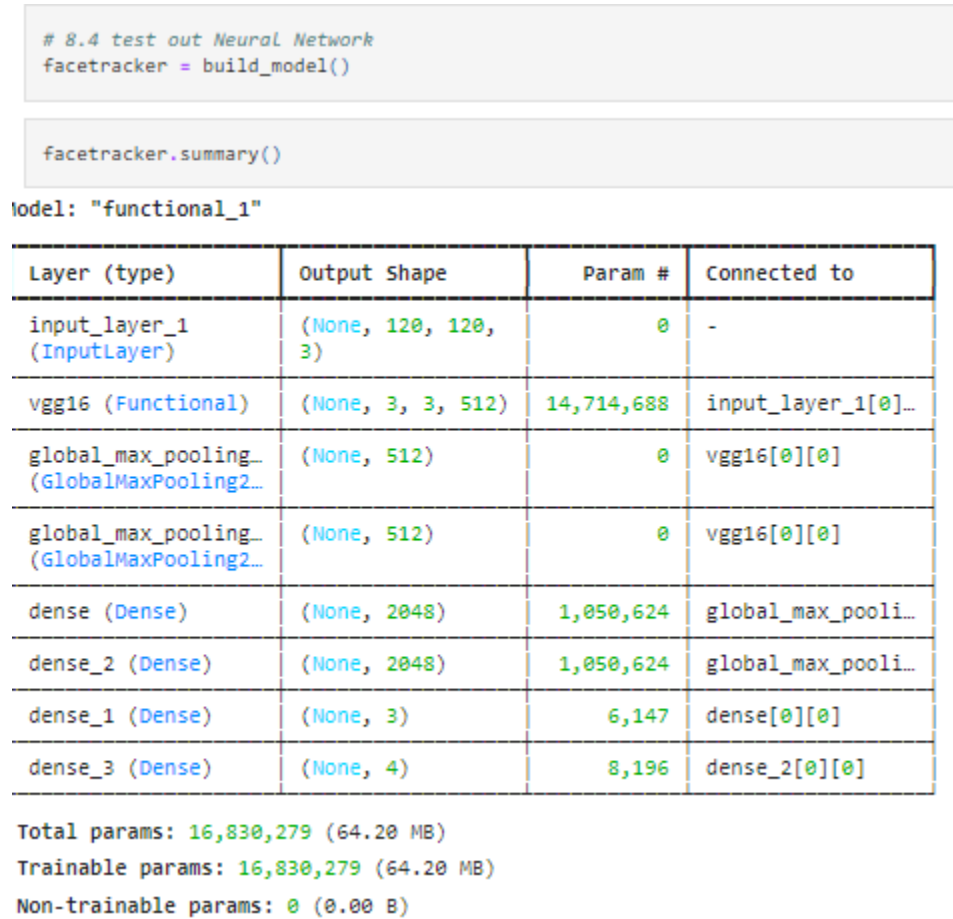


Figure 13: Loop to Draw Bounding Boxes on Images

Loading and Predicting with the Model

- Loading Data: Utilizing `train.as_numpy_iterator().next()` retrieves the next data batch from the training dataset. X is assigned for the input pictures, whereas y is designated for the labels and coordinates in this group. X contains the input images for the model, and y contains the classification labels and bounding box coordinates for the ground truth.

- Prediction: By using the pre-set facetracker model, the categories and positions of images in X are anticipated. The expected results are stored in categories and coordinates accordingly.

```
X, y = train.as_numpy_iterator().next()
```

```
classes, coords = facetracker.predict(X)
```

```
1/1 ----- 1s 550ms/step
```

Figure 14: Loading and Predicting with the Model

Localization Loss Function

- Coordinate Difference: This part calculates the squared difference between the real and predicted positions of the top-left corners of the bounding boxes. $y_true[:,2]$ shows the actual x and y coordinates of the true bounding boxes, while $yhat[:,2]$ shows the forecasted x and y coordinates of the bounding boxes.
- Size Difference: This indicates the difference in both the vertical and horizontal measurements of the real and predicted bounding boxes. The heights (h_true , h_pred) and widths (w_true , w_pred) are determined by subtracting the respective coordinates from each other.
- Total Loss: The localization loss is determined by summing the squared differences in coordinates and size, showing the extent of difference between the predicted and true bounding boxes.

```
# Creating Localization Loss and classification Loss

def localization_loss(y_true, yhat):
    delta_coord = tf.reduce_sum(tf.square(y_true[:, :2] - yhat[:, :2]))
    h_true = y_true[:, 3] - y_true[:, 1]
    w_true = y_true[:, 2] - y_true[:, 0]
    h_pred = yhat[:, 3] - yhat[:, 1]
    w_pred = yhat[:, 2] - yhat[:, 0]
    delta_size = tf.reduce_sum(tf.square(w_true - w_pred) + tf.square(h_true - h_pred))
    return delta_coord + delta_size
```

Figure 15: Localization loss and classification loss.

Classification of Loss Function

- Classification loss: ‘SparseCategoricalCrossentropy’ loss function is perfect for situations involving multi-class classification and integer labels.
- Regression loss: It is assigned to the designated localization_loss function discussed before, which calculates the discrepancy in determining bounding boxes.

```
classloss = tf.keras.losses.SparseCategoricalCrossentropy()
regressloss = localization_loss
```

```
# Testing out loss metrics
regressloss(y[1], coords)
```

```
<tf.Tensor: shape=(), dtype=float32, numpy=4.8303847>
```

```
classloss(y[0], classes)
```

```
<tf.Tensor: shape=(), dtype=float32, numpy=1.1383685>
```

Figure 16: Classifying the loss function.

Defining the FaceTracker Model

Initialization and Compilation

- Initialization: The FaceTracker class is a custom model that builds upon TensorFlow's Model class. It begins by assigning the eyetracker model to self.model.
- Compile Method: This technique sets up the model's optimizer, classification loss, and localization loss. It additionally calls the compile procedure from the parent class Model.

```
class FaceTracker(Model):  
    def __init__(self, eyetracker, **kwargs):  
        super().__init__(**kwargs)  
        self.model = eyetracker  
  
    def compile(self, opt, classloss, localizationloss, **kwargs):  
        super().compile(**kwargs)  
        self.closs = classloss  
        self.lloss = localizationloss  
        self.opt = opt
```

Figure 17: Defining the FaceTracker Model

Training Step

- Gradient Tape: It allows for the monitoring of operations to enable the calculation of gradients through automatic differentiation.
- Loss Calculation: The try block involves the model utilizing X to anticipate classes and coordinates. The original class labels (y[0]) are adjusted to fit, and the model computes both classification and localization losses. The total loss is the result of adding together these separate losses, each with varying degrees of importance.
- Error Handling: If a tf.errors.InvalidArgumentError occurs due to a zero batch size, a zero dummy loss will be given.
- Gradient Application: The optimizer computes and implements the gradients of the overall loss in relation to the trainable variables of the model.


```
@tf.function # Decorate train_step with @tf.function
def train_step(self, batch, **kwargs):
    X, y = batch

    with tf.GradientTape() as tape:
        try:
            classes, coords = self.model(X, training=True)

            # Ensure y[0] has a defined rank (handle potential reshaping)
            y_0 = tf.reshape(y[0], [-1, 1]) # Example: Reshape to (batch_size, 1) if needed

            # Check your data format and reshape accordingly
            # Ensure classes has a defined rank (check model output shape)
            # ... (reshape classes if necessary based on your model's output)

            batch_classloss = self.closs(y_0, classes)
            batch_localizationloss = self.lloss(tf.cast(y[1], tf.float32), coords)
            total_loss = batch_localizationloss + 0.5 * batch_classloss

        except tf.errors.InvalidArgumentError as e:
            # Handle the case where y[0] might have a batch size of zero (optional)
            if 'Input tensors must be of size at least 1' in str(e):
                return {"total_loss": tf.constant(0.0)} # Dummy Loss (optional)
            else:
                raise e # Re-raise other errors

        grad = tape.gradient(total_loss, self.model.trainable_variables)
        self.opt.apply_gradients(zip(grad, self.model.trainable_variables))

    return {"total_loss": total_loss, "class_loss": batch_classloss, "regress_loss": batch_localizationloss}
```

Figure 18: Training step

Testing Step

- Prediction: The model forecasts classes and coordinates of X without the need for training (training=False).
- Loss calculation involves transforming the actual class labels (y[0]) and then determining the losses for classification and localization. The total loss is calculated the same as it was during the training period.

```
def test_step(self, batch, **kwargs):
    X, y = batch
    classes, coords = self.model(X, training=False)

    # Ensure y[0] has a defined rank (handle potential reshaping)
    y_0 = tf.reshape(y[0], [-1, 1]) # Example: Reshape to (batch_size, 1) if needed
    # Check your data format and reshape accordingly
    # Ensure classes has a defined rank (check model output shape)
    # ... (reshape classes if necessary based on your model's output)
    batch_classloss = self.closs(y_0, classes)
    batch_localizationloss = self.lloss(tf.cast(y[1], tf.float32), coords)
    total_loss = batch_localizationloss + 0.5 * batch_classloss
    return {"total_loss": total_loss, "class_loss": batch_classloss, "regress_loss": batch_localizationloss}
```

Figure 18: Testing step

Compiling and Training the Model

- Model Initialization: Model is initialized by creating a FaceTracker object through the facetracker model.
- Putting together the model requires utilizing the optimizer, classification loss, and localization loss.
- Creating a folder for logs and setting up a TensorBoard callback to track training advancements.
- A single training epoch is carried out utilizing both the training and validation data, with the TensorBoard callback functionality activated.

```
model = FaceTracker(facetracker)

model.compile(opt, classloss, regressloss)

# defining Log folder
logdir = 'logs'
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=logdir)

hist = model.fit(train, epochs=1, validation_data=val, callbacks=[tensorboard_callback])

028/1028 ----- 1787s 2s/step - class_loss: 0.0866 - regress_loss: 0.1363 - total_loss: 0.1796 - val_class_loss: 0.1957 - val_regress_loss: 0.1365 - val_total_loss: 0.2344

# 10.3 Plot Performance
hist.history

{'class_loss': [0.0007510303985327482],
 'regress_loss': [0.010231394320726395],
 'total_loss': [0.010606909170746803],
 'val_class_loss': [0.19573935866355896],
 'val_regress_loss': [0.13653115928173065],
 'val_total_loss': [0.23440083861351013]}
```

Figure 19: Compiling and Training the Model

Making Predictions on the Test Dataset

- Test Data: Fresh batch of testing data has been acquired.

- Predicting: The facetracker model predicts the classes and coordinates of the test images.

```
# Making predictions on Test dataset  
test_data = test.as_numpy_iterator()
```

```
test_sample = test_data.next()
```

```
yhat = facetracker.predict(test_sample[0])
```

Figure 20: Predicting on test dataset.

Results

```
# drawing results
fig, ax = plt.subplots(ncols=8, figsize=(20,20))

for idx in range(8):
    className = ""
    sample_image = test_sample[0][idx].copy()

    classNumber = np.argmax(yhat[0][idx])

    if yhat[0][idx][classNumber] > 0.9:

        if classNumber > 0:
            sample_coord = yhat[1][idx]

            if classNumber == 1:
                className = "Eduardo"
            elif classNumber == 2:
                className = "Leo"

            cv2.rectangle(sample_image,
                          tuple(np.multiply(sample_coord[:2], [120,120]).astype(int)),
                          tuple(np.multiply(sample_coord[2:], [120,120]).astype(int)),
                          (255,0,0), 2)

    ax[idx].set_title(className)
    ax[idx].imshow(sample_image)
```

lipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
lipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
lipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
lipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
lipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
lipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



Figure 20: Result

Next Steps

- Work on statistical reports for IoU estimation models

Define the IoU calculation method: Clearly explain the IoU formula and approach.

Gather initial data: Create a dataset and calculate IoU values for each hypothesis.

Analyze IoU Distribution: Displays the IoU distribution and compares it to the baseline and provides a statistical summary.

- Evaluate the robustness of the model using images taken from different phones

Data Collection: Collect photos of the same article taken by different phone cameras.

Preprocessing: Standardize preprocessing steps such as scaling and normalization.

Model performance evaluation: Run the model on a new image, calculate the IoU score and analyze the robustness.

- Check the options to add more images and more complex models

Expand Dataset: Increase the size of the database by adding more resources and data.

Comparison: Evaluate the performance and feasibility of complex models.

Faced Challenges

- **Data Preparation and Handling:** One of the basic challenges was managing and arranging the data. Supervising different bunches (pictures and JSON names) and ensuring they were precisely stacked and preprocessed required cautious thought. The work for stacking pictures and names had to be critically sketched out to handle potential issues, such as record orchestrated inconsistencies and misplaced data. Besides, the resizing and normalization of pictures had to be standardized to ensure consistency over the dataset.
- **Balancing Dataset:** Ensuring that the dataset was balanced for planning was another challenge. The data augmentation process handle pointed to alter the dataset by making additional pictures and comparing names. In any case, affirming that the expanded data held the characteristics of the primary data and did not show bias was significant. The modifying, bunching, and prefetching of data to optimize planning execution as well required cautious tuning.
- **Model designing and Training:** Arranging a solid appear utilizing the functional API of TensorFlow and ensuring it facilitates well with VGG16 to highlight extraction was complex. The creation of divided branches for classification and backslide included to the complexity. Tuning the model's plan, selecting fitting sanctioning capacities, and setting the correct optimizer and learning rate were principal for fulfilling awesome execution.
- **Loss Function:** Executing custom loss function capacities for the localization of bounding boxes and combining it with classification mishap was an essential challenge. Ensuring that these loss functions precisely computed the contrasts and a deep understanding of the mathematical underpinnings and TensorFlow operations.

- Performance evaluation: Evaluating the model's execution on training and test datasets and ensuring that the given critical encounters were fundamental. Visualizing the comes almost, such as drawing bounding boxes on pictures and affirming the expected classes, was essential for subjective evaluation.
- Model Saving and Deployment: Saving the demonstration in an appropriate orchestrate that will be easily stacked for future acceptance or help planning was another bounce. Ensuring that the saved appear held all imperative components, checking weights, designing, and compilation state, was pivotal for reliable sending.

Lesson Learned

- **Comprehensive Data Handling:** Effective data managing with, and preprocessing are fundamental to the triumph of any machine learning amplify. Standardizing data bunches, ensuring consistency, and actualizing incredible data extension strategies basically influence illustrate execution. This wander braced the importance of data course of action and the requirement for seriously endorsement at each step.
- **Model Design Complexity:** Arranging models utilizing TensorFlow's utilitarian API offers uncommon flexibility but as well comes with complexity. This wander highlighted the noteworthiness of understanding the essential plan and instinctive between layers and components. A disengaged arranged and clear segment of assignments (classification and backslide) can lead to more practical and sensible models.
- **Importance of Custom Loss Functions:** Custom incident capacities custom fitted to specific errands, such as localization of bounding boxes, are significant for fulfilling needed execution. This wander ionized the requirement for a solid get a handle on of numerical concepts and TensorFlow operations to actualize and facilitate these custom capacities effectively.
- **Training Dynamics:** Watching and tuning the planning plan is fundamental to addressing issues such as overfitting, underfitting, and insecure points. This wander outlined the noteworthiness of utilizing methods like improving, clustering, prefetching, and lively learning rate modifications to protect planning soundness and make strides appear execution.

- Performance Metrics and Visualization: Evaluating appear execution utilizing appropriate estimations and visualizations gives imperative bits of information. This amplify underscored the require for both quantitative estimations and subjective examinations (such as visualizing bounding boxes) to urge it appear behavior comprehensively.
- Model Persistence and Deployment: Ensuring that models are saved precisely and can be easily stacked for course of action is critical.

Conclusion

Developing and teaching a complex neural network for identifying and verifying faces provided me with important practical knowledge in this significant project. The journey from data preparation to model design, training, evaluation, and saving presented numerous opportunities for learning and overcoming challenges. Crucial aspects to keep in mind include the vital importance of robust data management, the difficulties of developing models with advanced structures, the requirement for personalized loss functions, and the critical monitoring and adjustment of the training process. Additionally, it was essential to evaluate the model's efficacy using both numerical measurements and visualization tools, and to correctly store the model for implementation in the project.

Overall, this project not only strengthened fundamental concepts but also introduced more sophisticated techniques and best practices in deep learning, offering a comprehensive learning experience for all team members.